

# CartPole RL Enhancement

A Deep Reinforcement Learning Approach with Curriculum  
Learning and Energy Consumption Optimization

Louis Gauthier

January 13, 2025

# Table of Contents

Project Overview

Environment Description

Steps Towards the Solution

Possible Extensions

Possible Shortcomings

Results and Analysis

Conclusion

# Project Overview

- ▶ **Objective:** Enhance the classic CartPole environment by introducing a third action ("Do Nothing") and develop a deep RL solution to manage the modified environment.
- ▶ **Key Features:**
  - ▶ Addition of a neutral action to the action space.
  - ▶ Implementation of curriculum learning to adjust reward weights dynamically.
  - ▶ Incorporation of temperature-scaled softmax for exploration-exploitation balance.
  - ▶ Logging and monitoring using TensorBoard.
  - ▶ Energy consumption as a reward component.
- ▶ **Outcome:** Achieved a high success rate with efficient energy usage after 100k training steps.

# Environment Description

## ► **Classic CartPole:**

- Two actions: Apply a force to the cart to push it left or right.
- Observation Space: Cart position, cart velocity, pole angle, pole angular velocity.
- Objective: Balance the pole by moving the cart. Episode ends if pole angle exceeds a threshold or cart moves out of bounds.




Figure: The CartPole Environments

# Environment Description

- ▶ **Modified CartPole:**

- ▶ **Third Action:** *Do Nothing* - No force applied to the cart.
- ▶ **Reward Structure:** Split into four equal components: Alive, Distance to Center, Pole Angle, and Energy Usage.
- ▶ **Curriculum Learning:** Adjusts reward weights based on training phases.




Figure: The CartPole Environments

# Steps Towards the Solution

## 1. Environment Modification:

- ▶ Added a third "Do Nothing" action to the action space.
- ▶ Adjusted the reward function to incorporate energy consumption.

## 2. Algorithm Development:

- ▶ Implemented a combination of Double DQN and Dueling DQN for stable learning.
- ▶ Used a target network to improve training stability.
- ▶ Utilized Prioritized Experience Replay to sample important transitions.
- ▶ Applied Gradient Clipping to prevent exploding gradients.

## 3. Enhancements:

- ▶ Introduced Curriculum Learning to adjust reward weights dynamically.
- ▶ Integrated Temperature-Scaled Softmax for exploration-exploitation balance.
- ▶ Tested Normalization for observations and rewards.

# Steps Towards the Solution




Figure: Catastrophic Forgetting occurring in the training process

# Steps Towards the Solution

## 4. Problems encountered:

- ▶ Catastrophic forgetting: The agent forgets previously learned behaviors.
  - ▶ Solution: Use non-deterministic actions and temperature scheduling.
- ▶ Reward function is not a good metric of performance during training.
  - ▶ Solution: Run evaluations every 2k steps to track performance.

## 5. Logging and Monitoring:

- ▶ Set up TensorBoard for real-time monitoring of training metrics.
- ▶ Logged individual sub-rewards and penalties for detailed analysis.
- ▶ Evaluation every 2k steps to track performance.

## 6. Checkpointing:

- ▶ Enabled resuming training from saved checkpoints.
- ▶ Handling crashes or interruptions during training.



# TensorBoard Metrics

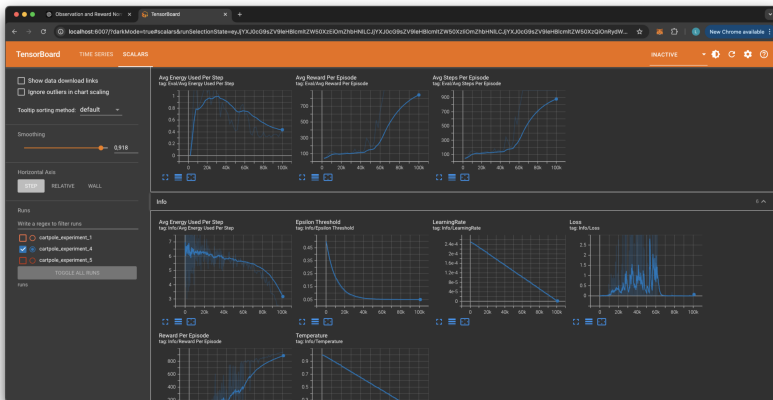


Figure: TensorBoard Dashboard showing training metrics

# Algorithm Enhancements

- ▶ **Double DQN:** Mitigates overestimation bias by decoupling action selection and evaluation.
- ▶ **Prioritized Experience Replay:** Samples important transitions more frequently to accelerate learning.
- ▶ **Normalization:**
  - ▶ *Observation Normalization:* Stabilizes training by normalizing input features.
  - ▶ *Reward Normalization:* Ensures consistent reward scaling.
- ▶ **Curriculum Learning:**
  - ▶ Phases:
    1. **0-20k steps:** 50% Alive Reward, 50% Pole Angle Reward.
    2. **20k-50k steps:** 33% Alive, 33% Pole Angle, 33% Distance to Center.
    3. **Above 50k steps:** 25% each for Alive, Pole Angle, Distance to Center, and Energy Usage.

# Algorithm Enhancements

## ► Temperature-Scaled Softmax:

- Initial Temperature: 1.0
- Final Temperature: 0.1
- Linearly decreases over training to reduce exploration over time.




Figure: Loss Stabilization: Prioritized (blue) vs. Uniform Replay (red)

# Possible Extensions

- ▶ **Stratification:**
  - ▶ Implement stratified sampling to ensure diverse experiences in replay buffer.
- ▶ **Distribution-Aware Buffer:**
  - ▶ Explore Cosine Similarity or Variational Autoencoders (VAEs) to model state distributions and enhance replay buffer sampling.
- ▶ **Parallel Environments:**
  - ▶ Increase to 16 parallel environments for faster experience collection and improved training efficiency.
- ▶ **Reward Function Refinement:**
  - ▶ Schedule Pole Angle and Distance to Center rewards to decrease over time, focusing more on energy optimization.
- ▶ **Model Checkpointing:**
  - ▶ Maintain multiple checkpoints to preserve best-performing models across different training phases.
- ▶ **Advanced Exploration Strategies:**
  - ▶ Incorporate other exploration methods like entropy regularization or intrinsic motivation (e.g., curiosity).

# Possible Shortcomings

- ▶ **Algorithm Tuning:**

- ▶ Deep RL algorithms require extensive hyperparameter tuning for optimal performance.

- ▶ **Overfitting:**

- ▶ The model may overfit to the specific environment configuration, limiting generalizability.

- ▶ **Fixed Length Episodes:**

- ▶ Fixed-length episodes combined with energy consumption reward may lead to suboptimal behavior where the agent overfits to short-term gains.

- ▶ **Training For Too Long Might Decrease Performance:**


- ▶ The agent might overfit to near-perfect behavior and forget more complex situations.

- ▶ **Learning Stability:**

- ▶ Other on-policy algorithm like PPO might be more stable and efficient.

# Results and Analysis

- ▶ **Training Outcome after 100k Steps:**
  - ▶ **Average Reward:** 975/1000
  - ▶ **Energy Consumption:** Under 0.4 per step (Full Power = 10 per step)
  - ▶ **Success Rate:** 100% during evaluation (No failures)
- ▶ **Evaluation Strategy:**
  - ▶ Conducted evaluations every 2k steps.
  - ▶ Essential for monitoring performance due to off-policy learning, non-deterministic actions, and exploration.
- ▶ **Prioritization Analysis:**
  - ▶ Prioritized Experience Replay helped stabilize loss.
  - ▶ Did not significantly enhance overall performance compared to uniform sampling.



# Conclusion

- ▶ Successfully enhanced the CartPole environment with a third "Do Nothing" action and optimized reward structure for energy consumption.
- ▶ Developed a robust deep RL solution incorporating curriculum learning and temperature-scaled softmax for effective exploration-exploitation balance.
- ▶ Achieved high performance with minimal energy usage and a perfect success rate in evaluations.
- ▶ Identified areas for improvement, including further algorithm tuning and exploring advanced buffer management techniques.